

# Sample Code for an Autoencoder for Dr. Goswami

**Shakil Rafi**

## Introduction and Dataset

In this notebook we will do a test-case with synthetic data comparing classic PCA with an autoencoder to see that the autoencoder is able to take high-dimensional data and extract out the relevant features.

## Dataset

Our dataset will be a synthetic dataset in which we will use the sklearn datasets method to create eight “blobs” of data and 30,000 observations. Each blob has a centroid around which the data is clustered in a Gaussian fashion with a standard deviation of 2, i.e. each element of the cluster  $c$ ,  $x_c$  is such that  $x_c \sim \mathcal{N}(\text{centroid}_c, c)$ . We will have nine features to the data, to which we will add an extra feature, consisting of noise coming from a uniform distribution.

*Note on dataset:* The author regrets to inform that due to the ongoing shutdowns of internet and communications in Bangladesh that they have had difficulty and delay obtaining actual datasets. This Jupyter notebook therefore serves as a proof of concept that autoencoders can be beneficial and viable in analyzing genomic data via proxy with synthetic data. The author apologizes for this inconvenience.

## Introduction and Methodology

Our methodology will be as follows:

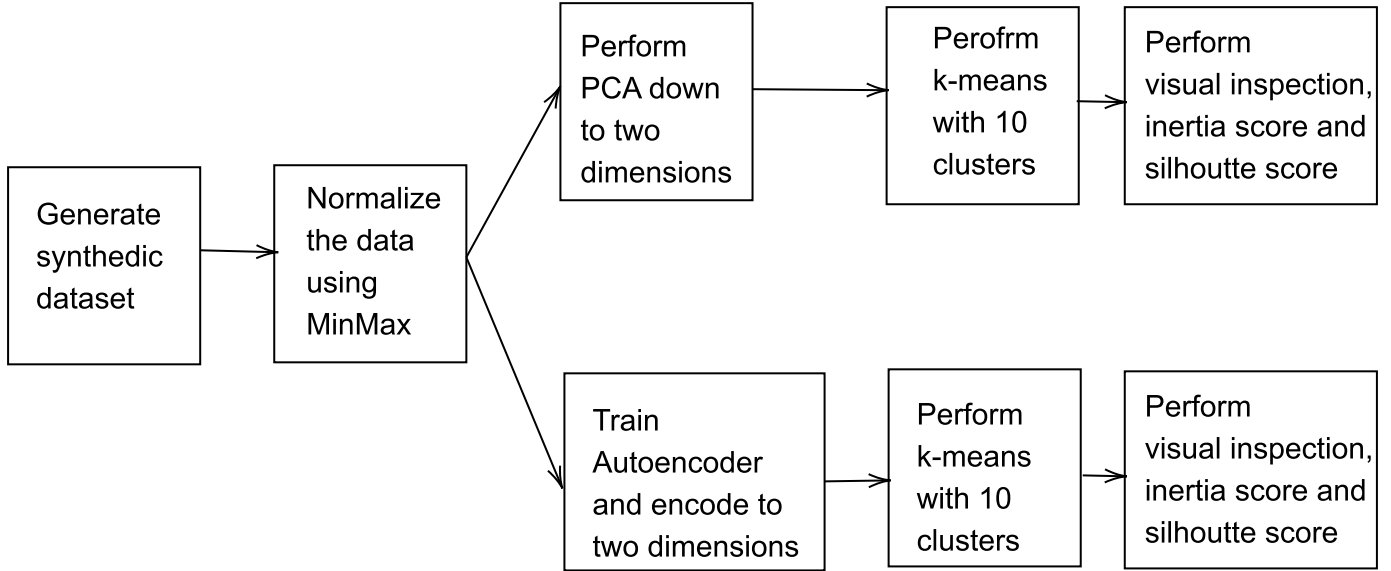


Figure 1: Methodology

In summary, we will take our synthetic dataset and perform normalization using the MinMax Scaler. Upon completion of the PCA transform we will perform a scatterplot and perform a basic visual inspection. Ideally, we would expect to see eight clusters, but even here we will see that PCA will fail dramatically. We will then perform a classic k-means clustering on this reduced space. We will then calculate the relevant scores inertia defined as:

$$\text{Inertia} = \sum_{i=1}^n \sum_{j=1}^k \|x_i - \mu_j\|^2 \cdot \mathbf{1}_{x_i \in C_j}$$

where  $n$  is the number of data points,  $k$  is the number of clusters  $\mu_j$  is the centroid of each cluster.

Essentially, the inertial score tells us how “compact” each cluster is. Scores range from  $[0, \infty)$ . A lower score means a better cluster, although this is not the only metric for measuring clustering effectiveness. One other measure is silhouette score, defined for a single datapoint  $x_i$  as:

$$\text{Silhouette}(x_i) = \frac{b(x_i) - a(x_i)}{\max\{a(x_i), b(x_i)\}}$$

Where  $a(x_i)$  is the average distance between  $x_i$  and all other points in the cluster, mathematically speaking, this is:

$$a(x_i) = \frac{1}{|C_i|-1} \left( \sum_{x_j \in C_k, x_j \neq x_i} \|x_i - x_j\| \right)$$

And where  $b(x_i)$  is the minimum average distance between  $x_i$  and all other points in other clusters, mathematically speaking, this is:

$$b(x_i) = \min_{C_m \neq C_k} \left( \frac{1}{|C_m|} \sum_{x_j \in C_m} \|x_i - x_j\| \right)$$

```
#####
#
# The current cell installs packages if not already installed in the hosts
# systems.
# This is supplemented by an extensive requirements.txt file in the folder
#
#####

# import subprocess
# import sys

# def install(package):
#     subprocess.check_call([sys.executable, "-m", "pip", "install", package])

# # List your packages here
# packages = ["numpy", "pandas", "seaborn", "keras", "tensorflow", "matplotlib",
# "scikit-learn"]

# for package in packages:
#     install(package)
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
from sklearn.datasets import make_blobs
```

```
data = make_blobs(
    n_samples=30000,
    n_features=9,
    centers=8,
    cluster_std=2
```

```
)

X,y = data
```

```
# Create a random uniform column of noise to be added as a feature

np.random.seed(seed=101)
noise = np.random.uniform(size=len(X))
noise = pd.Series(noise)
```

```
# Add noise to our dataset. Note that since no one dimension is any
# special than another dimension it does not matter where we add
# the noise to

feat = pd.DataFrame(X)
feat = pd.concat([feat,noise],axis=1)
```

```
# Rename the columns to be human readable

feat.columns = [f"X{i+1}" for i in range(len(feat.columns))]
```

```
# Visually inspect our X

feat
```

	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10
0	3.832150	-0.885521	1.396813	6.518364	-7.181659	7.624980	9.863516	-3.958961	10.142876	0.516399
1	5.121997	-9.531866	6.580018	3.576645	-3.133651	3.158619	7.563490	6.267551	-0.000102	0.570668
2	0.170303	-3.953752	3.746255	4.652394	1.368011	-0.426606	2.571776	-7.456172	9.752707	0.028474
3	2.087428	-2.992573	5.104202	6.203770	0.087370	-4.681826	5.966582	-0.207810	7.504352	0.171522
4	9.513307	6.389279	-2.908372	0.090795	6.912397	2.130114	-9.692870	5.897349	1.681667	0.685277
...	...	...	...	...	...	...	...	...	...	...
29995	13.124232	8.117214	8.370973	9.811386	0.173668	4.734505	2.027921	-5.618596	7.445326	0.655720
29996	2.451051	-4.309338	3.712033	5.494899	0.206910	-1.965390	4.748489	-4.381654	7.359769	0.298044
29997	-0.500273	0.921920	2.131400	3.418651	-6.197750	0.141882	5.531861	2.188421	5.244825	0.537972
29998	3.285187	0.026805	6.310630	5.544782	-0.692804	0.721316	4.952729	-3.549326	3.656933	0.513805
29999	10.510771	-6.328579	12.581769	9.529466	-3.849963	4.845036	3.824897	-1.863592	5.485848	0.509921

```
# Summary statistics of our dataset for EDA
```

```
feat.describe()
```


	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	6.880943	1.773305	-0.392525	5.431513	-2.360405	1.193830	1.046358	-0.117281	-0.424195	0.500133
std	3.450987	7.130533	7.150479	3.050602	4.372199	5.664284	6.621440	4.890498	5.874961	10.289540
min	-5.490365	17.244616	17.053376	64.412489	17.099587	13.346374	16.299739	12.580435	16.440138	0.000002
25%	4.730572	4.391468	7.417903	3.151284	-5.552574	3.265249	5.086935	3.982851	4.665434	0.247814
50%	7.459597	3.072980	0.064210	5.355945	-1.532697	0.390769	2.553161	-1.276105	1.152058	0.499429
75%	9.374422	8.119963	6.353483	7.721868	1.037876	6.500394	6.235492	3.957791	3.830021	0.750924
max	17.518416	6.887564	4.705751	15.562946	0.309994	15.120596	6.109644	4.901781	15.729132	0.999985

```
# Save our dataset
```

```
feat.to_csv("feat.csv")
```

```
# 3D scatterplot of three arbitrary dimensions
```

```
fig = plt.figure()
ax = fig.add_subplot(111, projection = '3d')
ax.scatter(feat['X1'], feat['X2'], feat['X9'], c = y)
```



### Preprocessing the data

We will use the standard MinMaxScaler from sklearn to scale and preprocess the data

```
# Preprocess our data using MinMaxScaler

from sklearn.preprocessing import MinMaxScaler

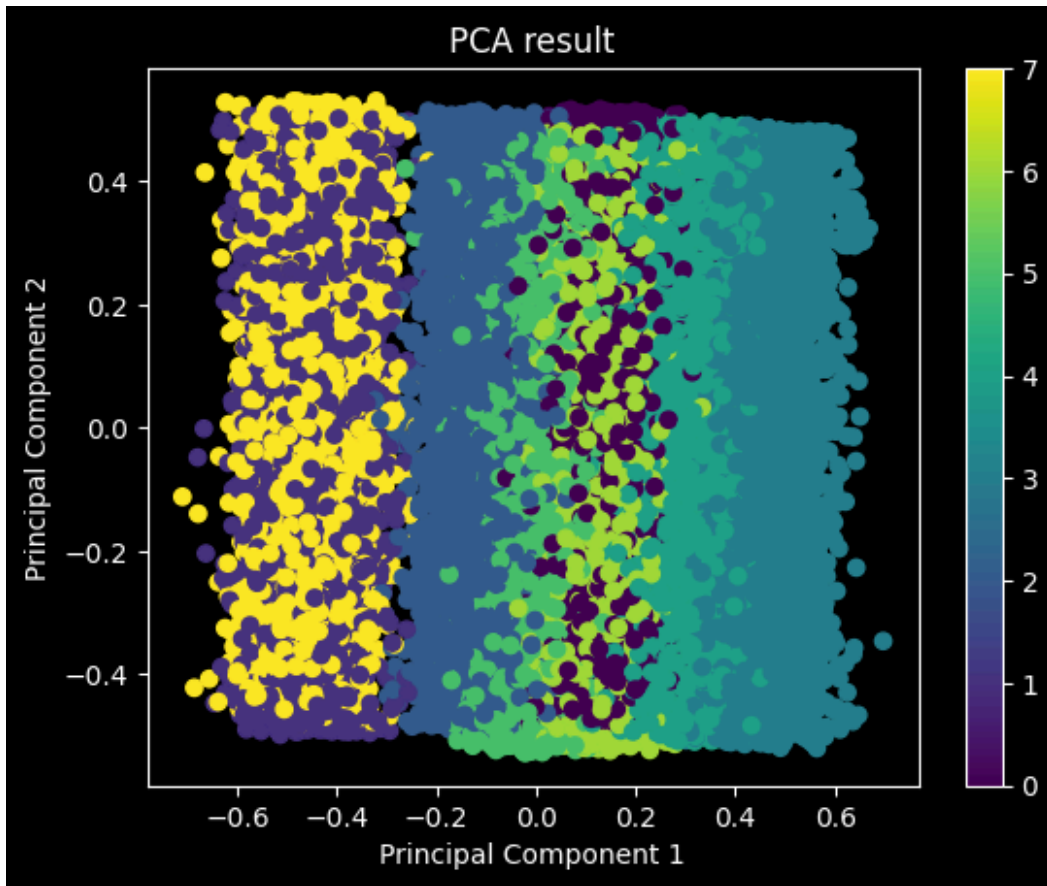
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(feats)
```

### The PCA decomposition

```
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
pca_result = pca.fit_transform(scaled_data)

plt.scatter(pca_result[:, 0], pca_result[:, 1], c=y)
plt.title('PCA result')
plt.colorbar()
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.show()
```



Note that PCA fails dramatically at this dimension reduction task. A priori we know that this dataset has clusters as it was constructed this way but PCA fails to capture the clustering. This is because PCA is a linear transformation and fails to capture any non-linear trends in the data. Note also that the uniform noise dimension results in a projection that is “smeared” yielding no obvious clusters on visual inspection. We will run two separate metrics later on to show how much more effective autoencoders are than regular PCA as a pre-step towards clustering.

```
from sklearn.metrics import mean_squared_error

X_reconstructed = pca.inverse_transform(pca_result)
print("Reconstruction loss of PCA:", mean_squared_error(X_reconstructed,
scaled_data))
```

```
Reconstruction loss of PCA: 0.021622192886071475
```

### K-means done on the encoded data

We now seek to apply classic K-Means on this PCA-result. We will do the same on the

```
from sklearn.cluster import KMeans
```

```
kmeans_for_pca = KMeans(n_clusters=8)  
kmeans_for_pca.fit(pca_result)
```

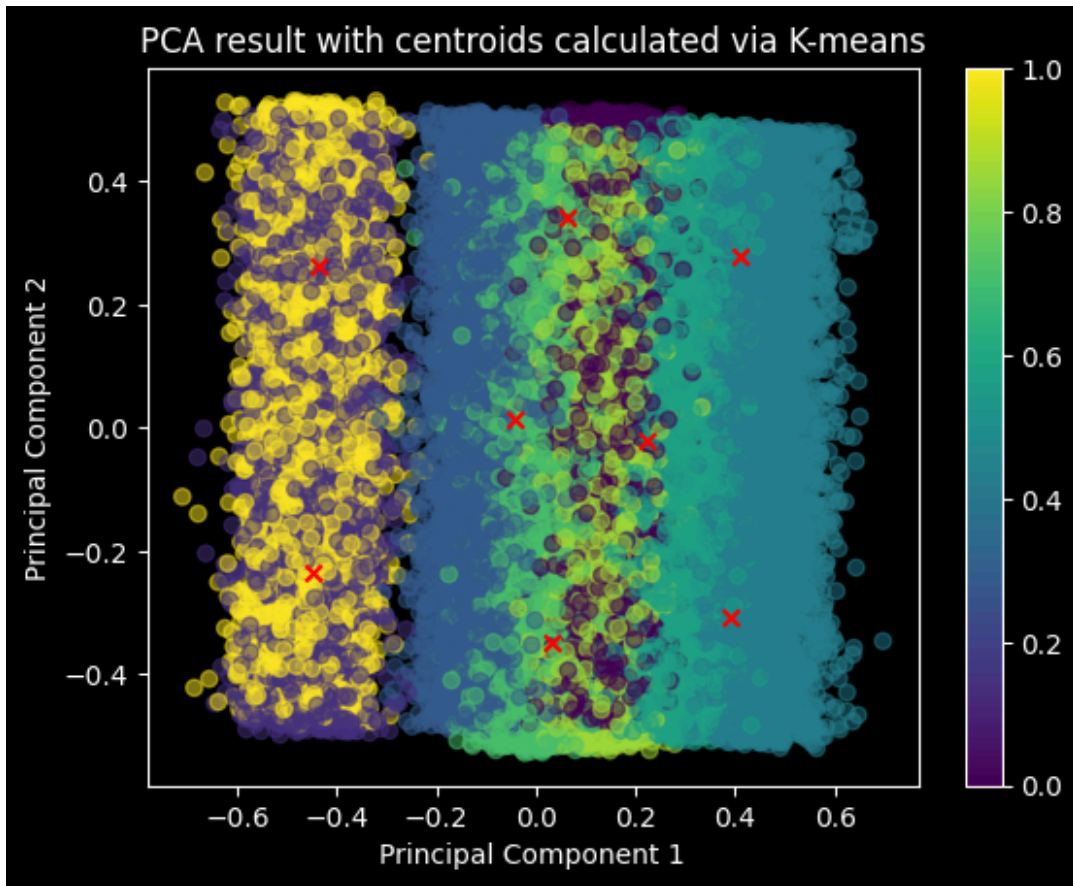
```
KMeans()
```

```
kmeans_for_pca.cluster_centers_
```

```
array([[ -0.44910219, -0.23529575],  
       [ -0.04385551,  0.0132341 ],  
       [ -0.43505372,  0.26351184],  
       [  0.03283007, -0.34681073],  
       [  0.39055912, -0.30722012],  
       [  0.40882084,  0.27793082],  
       [  0.22417627, -0.01915176],  
       [  0.06386291,  0.34105549]])
```

```
plt.scatter(pca_result[:, 0], pca_result[:, 1], c=y, alpha=0.5)  
plt.scatter(kmeans_for_pca.cluster_centers_[0],  
            kmeans_for_pca.cluster_centers_[0,1], marker='x', c='red')  
plt.title('PCA result with centroids calculated via K-means')  
plt.colorbar()  
plt.xlabel("Principal Component 1")  
plt.ylabel("Principal Component 2")  
plt.show()
```





```
kmeans_for_pca.inertia_
```

```
665.493846052085
```

### The autoencoder decomposition

We will now create an autoencoder. We will choose a rather deep model, with  $10 \rightarrow 5 \rightarrow 2 \rightarrow 5 \rightarrow 10$  layer widths. Our autoencoder will first encode the 10 dimensional data to a 2 dimensional latent space, and then decode it back to 10 dimensions. It will be trained against our dataset until it learns an optimal representation for this dataset into two latent dimensions.

We will then extract out the encoder and visually inspect our 2d latent space. Because neural networks are efficient at learning the underlying structure of unstructured data it's latent space should be noise-free, moreso than the PCA representation. As such, even visually we should see much more clustering than with PCA. Schematically, our autoencoder looks as follows:




Figure 2: Autoencoder

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.optimizers import SGD
```

```
# # This is the input layer
# input_layer = Input(shape=(scaled_data.shape[1],))

# # This is the encoding layers
# encoded = Dense(8, activation='relu')(input_layer)
# encoded = Dense(6, activation='relu')(encoded)
# encoded = Dense(2, activation='relu')(encoded)

# # This is the decoding layers
# decoded = Dense(2, activation='relu')(encoded)
# decoded = Dense(6, activation='relu')(decoded)
# decoded = Dense(8, activation='relu')(decoded)
# decoded = Dense(scaled_data.shape[1], activation='linear')(decoded)

# # This is the full autoencoder
# autoencoder = Model(inputs=input_layer, outputs=decoded)
```

```
input_layer = Input(shape=(scaled_data.shape[1],))

# Define the encoding layers
encoded = Dense(8, activation='sigmoid')(input_layer)
encoded = Dense(6, activation='sigmoid')(encoded)
encoded_output = Dense(2, activation='linear')(encoded)
```

```

# Define the encoder model
encoder = Model(inputs=input_layer, outputs=encoded_output)

# Define the decoding layers
encoded_input = Input(shape=(2,))
decoded = Dense(6, activation='sigmoid')(encoded_input)
decoded = Dense(8, activation='sigmoid')(decoded)
decoded_output = Dense(scaled_data.shape[1], activation='sigmoid')(decoded)

# Define the decoder model
decoder = Model(inputs=encoded_input, outputs=decoded_output)

# Define the autoencoder model
autoencoder_input = Input(shape=(scaled_data.shape[1],))
encoded_repr = encoder(autoencoder_input)
reconstructed = decoder(encoded_repr)

autoencoder = Model(inputs=autoencoder_input, outputs=reconstructed)

# Print model summaries
print("Encoder summary:")
encoder.summary()
print("\nDecoder summary:")
decoder.summary()
print("\nAutoencoder summary:")
autoencoder.summary()

```

Encoder summary:

Decoder summary:

Autoencoder summary:

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 10)	0
dense (Dense)	(None, 8)	88
dense_1 (Dense)	(None, 6)	54

dense_2 (Dense)	(None, 2)	14
-----------------	-----------	----

Total params: 156 (624.00 B)

Trainable params: 156 (624.00 B)

Non-trainable params: 0 (0.00 B)

Model: "functional\_1"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 2)	0
dense_3 (Dense)	(None, 6)	18
dense_4 (Dense)	(None, 8)	56
dense_5 (Dense)	(None, 10)	90

Total params: 164 (656.00 B)

Trainable params: 164 (656.00 B)

Non-trainable params: 0 (0.00 B)

Model: "functional\_2"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 10)	0
functional (Functional)	(None, 2)	156

functional_1 (Functional)	(None, 10)	164
---------------------------	------------	-----

Total params: 320 (1.25 KB)

Trainable params: 320 (1.25 KB)

Non-trainable params: 0 (0.00 B)

```
autoencoder.summary()
```

Model: "functional\_2"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 10)	0
functional (Functional)	(None, 2)	156
functional_1 (Functional)	(None, 10)	164

Total params: 320 (1.25 KB)

Trainable params: 320 (1.25 KB)

Non-trainable params: 0 (0.00 B)

```
autoencoder.compile(loss='mse',
                    optimizer=SGD(learning_rate=0.01))
```

```
autoencoder.fit(scaled_data,
                scaled_data,
                batch_size = 64,
                shuffle = True,
                epochs = 50)
```

```

Epoch 1/50
469/469 _____ 0s 391us/step - loss: 0.0658
Epoch 2/50
469/469 _____ 0s 793us/step - loss: 0.0600
Epoch 3/50
469/469 _____ 0s 381us/step - loss: 0.0549
Epoch 4/50
469/469 _____ 0s 355us/step - loss: 0.0507
Epoch 5/50
469/469 _____ 0s 374us/step - loss: 0.0478
Epoch 6/50
469/469 _____ 0s 372us/step - loss: 0.0453
Epoch 7/50
469/469 _____ 0s 336us/step - loss: 0.0436
Epoch 8/50
469/469 _____ 0s 340us/step - loss: 0.0423
Epoch 9/50
469/469 _____ 0s 356us/step - loss: 0.0414
Epoch 10/50
469/469 _____ 0s 356us/step - loss: 0.0409
Epoch 11/50
469/469 _____ 0s 354us/step - loss: 0.0404
Epoch 12/50
469/469 _____ 0s 348us/step - loss: 0.0400
Epoch 13/50
469/469 _____ 0s 352us/step - loss: 0.0399
Epoch 14/50
469/469 _____ 0s 350us/step - loss: 0.0398
Epoch 15/50
469/469 _____ 0s 348us/step - loss: 0.0398
Epoch 16/50
469/469 _____ 0s 349us/step - loss: 0.0397
Epoch 17/50
469/469 _____ 0s 429us/step - loss: 0.0397
Epoch 18/50
469/469 _____ 0s 380us/step - loss: 0.0396
Epoch 19/50
469/469 _____ 0s 501us/step - loss: 0.0397
Epoch 20/50
469/469 _____ 0s 369us/step - loss: 0.0395
Epoch 21/50
469/469 _____ 0s 388us/step - loss: 0.0396
Epoch 22/50
469/469 _____ 0s 367us/step - loss: 0.0396
Epoch 23/50
469/469 _____ 0s 361us/step - loss: 0.0396
Epoch 24/50
469/469 _____ 0s 345us/step - loss: 0.0395

```

```

Epoch 25/50
469/469 _____ 0s 344us/step - loss: 0.0396
Epoch 26/50
469/469 _____ 0s 383us/step - loss: 0.0396
Epoch 27/50
469/469 _____ 0s 334us/step - loss: 0.0397
Epoch 28/50
469/469 _____ 0s 342us/step - loss: 0.0397
Epoch 29/50
469/469 _____ 0s 340us/step - loss: 0.0396
Epoch 30/50
469/469 _____ 0s 341us/step - loss: 0.0394
Epoch 31/50
469/469 _____ 0s 339us/step - loss: 0.0396
Epoch 32/50
469/469 _____ 0s 348us/step - loss: 0.0396
Epoch 33/50
469/469 _____ 0s 348us/step - loss: 0.0396
Epoch 34/50
469/469 _____ 0s 339us/step - loss: 0.0396
Epoch 35/50
469/469 _____ 0s 340us/step - loss: 0.0396
Epoch 36/50
469/469 _____ 0s 342us/step - loss: 0.0396
Epoch 37/50
469/469 _____ 0s 641us/step - loss: 0.0395
Epoch 38/50
469/469 _____ 0s 342us/step - loss: 0.0395
Epoch 39/50
469/469 _____ 0s 341us/step - loss: 0.0395
Epoch 40/50
469/469 _____ 0s 345us/step - loss: 0.0396
Epoch 41/50
469/469 _____ 0s 344us/step - loss: 0.0395
Epoch 42/50
469/469 _____ 0s 347us/step - loss: 0.0397
Epoch 43/50
469/469 _____ 0s 374us/step - loss: 0.0396
Epoch 44/50
469/469 _____ 0s 343us/step - loss: 0.0395
Epoch 45/50
469/469 _____ 0s 340us/step - loss: 0.0396
Epoch 46/50
469/469 _____ 0s 341us/step - loss: 0.0396
Epoch 47/50
469/469 _____ 0s 337us/step - loss: 0.0395
Epoch 48/50
469/469 _____ 0s 345us/step - loss: 0.0395

```


```
Epoch 49/50  
469/469 _____ 0s 342us/step - loss: 0.0396  
Epoch 50/50  
469/469 _____ 0s 348us/step - loss: 0.0395
```

```
<keras.src.callbacks.history.History at 0x30913c2c0>
```

```
# Encode our data down to two dimensions using our  
# trained autoencoder  
  
encoder = Model(inputs=input_layer, outputs=encoded)  
encoded_2dim = encoder.predict(scaled_data)
```

```
938/938 _____ 1s 517us/step
```

```
plt.scatter(encoded_2dim[:,0],encoded_2dim[:,1], c = y)
```



Visually we see that our autoencoder has performed much better than the PCA. We see four clusters very clearly. Not only that the clusters are much more compact than with PCA.

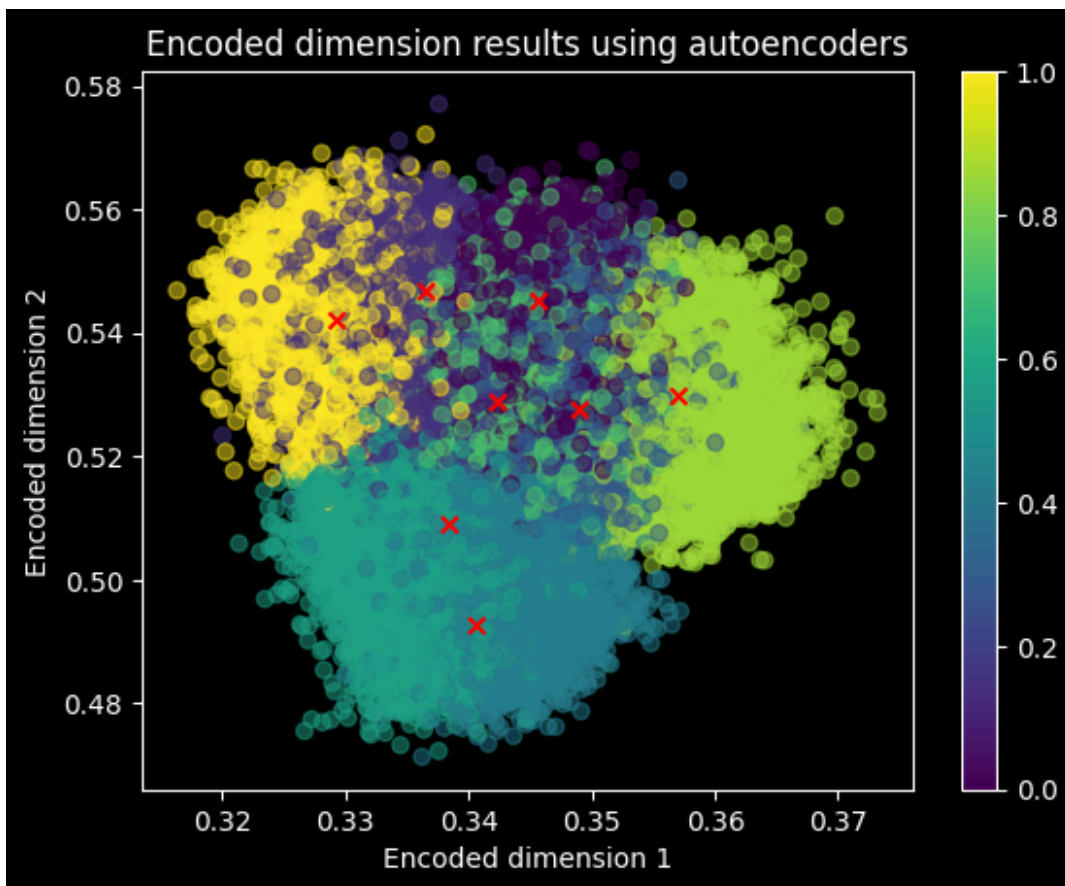


```
# Train our k-means on the encoded 2 dimensional data
```

```
kmeans_for_autoencoder = KMeans(n_clusters=8)  
kmeans_for_autoencoder.fit(encoded_2dim)
```

```
KMeans()
```

```
plt.scatter(encoded_2dim[:, 0], encoded_2dim[:, 1], c=y, alpha=0.5)  
plt.scatter(kmeans_for_autoencoder.cluster_centers_[0],  
            kmeans_for_autoencoder.cluster_centers_[1], marker='x', c='red')  
plt.title('Encoded dimension results using autoencoders')  
plt.colorbar()  
plt.xlabel("Encoded dimension 1")  
plt.ylabel("Encoded dimension 2")  
plt.show()
```



```
kmeans_for_autoencoder.inertia_
```

```
6.157865047454834
```

```
kmeans_for_pca.inertia_
```

```
665.493846052085
```

```
kmeans_for_pca.labels_
```

```
array([6, 2, 3, ..., 6, 1, 2], dtype=int32)
```

```
from sklearn.metrics import silhouette_score
```

```
silhouette_score(scaled_data, kmeans_for_pca.labels_)
```

```
0.13132647285236687
```

```
silhouette_score(scaled_data, kmeans_for_autoencoder.labels_)
```

```
0.19285495491413215
```

## Saving our models

Finally we will save our model for future use

```
autoencoder.save("autoencoder.keras")
```

```
import pickle

with open('pca_model.pkl', 'wb') as file:
    pickle.dump(pca, file)
```